

# Unikernels - The next big little thing ?

Devconf.cz, Brno, Czech Republic, Jan 2017



Mike Bright,  @mjbright

# Unikernels

## Overview

- What are Unikernels?
  - Why do we need them?
  - Domains of application
- Unikernel implementations
  - Clean-slate or legacy
  - Tooling
- Containers and Unikernels
- Demo
- Conclusions

# Why this talk?

@mjbright

# Why this talk?

Curiosity about

- What we can expect to see from Unikernels (and Docker ...)
- Who the players are

# Why this talk?

Curiosity about

- What we can expect to see from Unikernels (and Docker ...)
- Who the players are

1990's: First unikernels - Exokernel and Nemesis (Univ. Cambr)

# Why this talk?

Curiosity about

- What we can expect to see from Unikernels (and Docker ...)
- Who the players are

1990's: First unikernels - Exokernel and Nemesis (Univ. Cambr)

# Why this talk?

Curiosity about

- What we can expect to see from Unikernels (and Docker ...)
- Who the players are

1990's: First unikernels - Exokernel and Nemesis (Univ. Cambr)

# Why this talk?

Curiosity about

- What we can expect to see from Unikernels (and Docker ...)
- Who the players are

1990's: First unikernels - Exokernel and Nemesis (Univ. Cambr)



# Why this talk?

Curiosity about

- What we can expect to see from Unikernels (and Docker ...)
- Who the players are

1990's: First unikernels - Exokernel and Nemesis (Univ. Cambr)

# Why this talk?

Curiosity about

- What we can expect to see from Unikernels (and Docker ...)
- Who the players are

1990's: First unikernels - Exokernel and Nemesis (Univ. Cambr)

# Why this talk?

Curiosity about

- What we can expect to see from Unikernels (and Docker ...)
- Who the players are

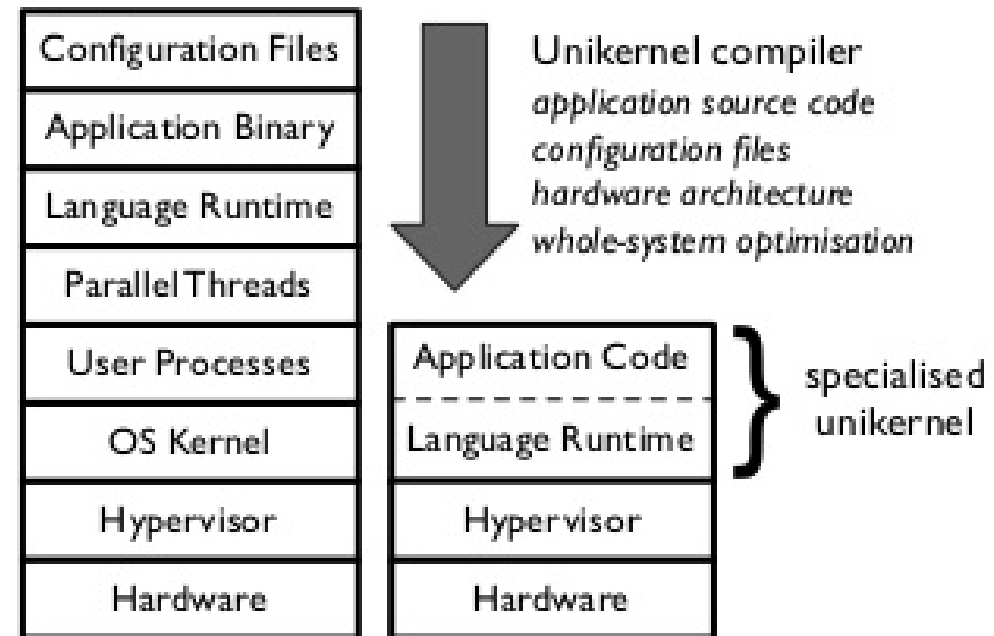
1990's: First unikernels - Exokernel and Nemesis (Univ. Cambr)

In 2017 ?

- MirageOS 3 will be released
- DockerCon US and EU

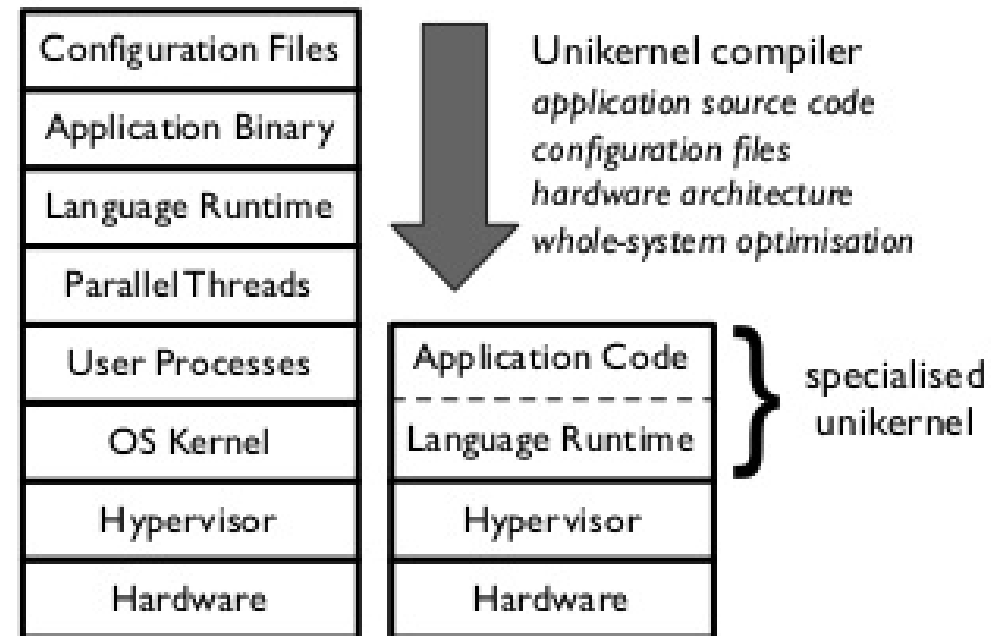
# What are Unikernels? "Library OS"

Applications built with only the OS components they actually require, e.g. TCP Stack, DNS, DHCP, NAT, F/w, Disk access.



# What are Unikernels? "Library OS"

Applications built with only the OS components they actually require, e.g. TCP Stack, DNS, DHCP, NAT, F/w, Disk access.



Single process(\*) applications (no threads, forking or multi-user) (\*)

Small size (few lines of code) and very fast to boot

Small attack surface (potentially secure)

@mjbright

High performance - no context switches!

No shell

# Why are Unikernels needed?

Think for a moment - What OS do you run?

On what hardware?

Modern OS provide amazing backwards compatibility and features

# Why are Unikernels needed?

Think for a moment - What OS do you run?

On what hardware?

Modern OS provide amazing backwards compatibility and features

But an app uses a tiny fraction of those features consuming resources,  
increasing the attack surface (linux kernel ~ 25 M LOC)


# Why are Unikernels needed?

Think for a moment - What OS do you run?

On what hardware?

Modern OS provide amazing backwards compatibility and features

But an app uses a tiny fraction of those features consuming resources,  
increasing the attack surface (linux kernel ~ 25 M LOC)





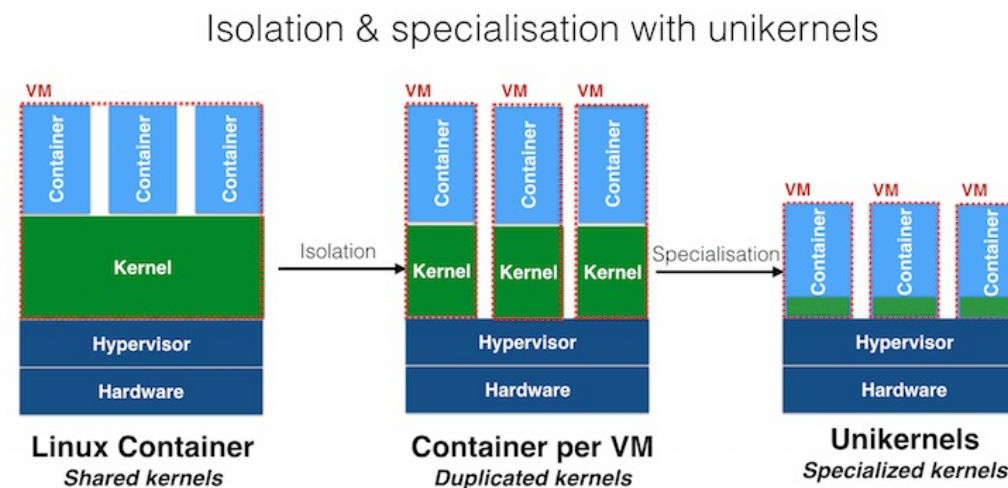
# Why are Unikernels needed?

Think for a moment - What OS do you run?

On what hardware?

Modern OS provide amazing backwards compatibility and features

But an app uses a tiny fraction of those features consuming resources, increasing the attack surface (linux kernel ~ 25 M LOC)



@mjbright

Unikernels provide an alternative

But are they a panacea?

# In what domains might they be used?

## Cloud Computing

- Small (kB/MB) immutable entities with fast boot times (100's ms).
- Possibility of on-demand servers,  $\mu$ -services
- Potentially greater security (< LOC)

# In what domains might they be used?

## Cloud Computing

- Small (kB/MB) immutable entities with fast boot times (100's ms).
- Possibility of on-demand servers,  $\mu$ -services
- Potentially greater security (< LOC)

## NFV (Network Functional Virtualization)

- As cloud but stricter requirements on response times, service chaining

 [Unikernels meet NFV](#); [Ericsson Research Blog](#); [Unikernels.org Blog](#)

# In what domains might they be used?

## Cloud Computing

- Small (kB/MB) immutable entities with fast boot times (100's ms).
- Possibility of on-demand servers,  $\mu$ -services
- Potentially greater security ( $<$  LOC)

## NFV (Network Functional Virtualization)

- As cloud but stricter requirements on response times, service chaining

 [Unikernels meet NFV](#); [Ericsson Research Blog](#); [Unikernels.org Blog](#)

## IoT / Embedded

- For low-resource, potentially secure elements (baremetal or  $\mu$ -vmm ?)
- Build up the "app" instead of stripping down the "OS"

# In what domains might they be used?

## Cloud Computing

- Small (kB/MB) immutable entities with fast boot times (100's ms).
- Possibility of on-demand servers,  $\mu$ -services
- Potentially greater security (< LOC)

## NFV (Network Functional Virtualization)

- As cloud but stricter requirements on response times, service chaining

 [Unikernels meet NFV](#); [Ericsson Research Blog](#); [Unikernels.org Blog](#)

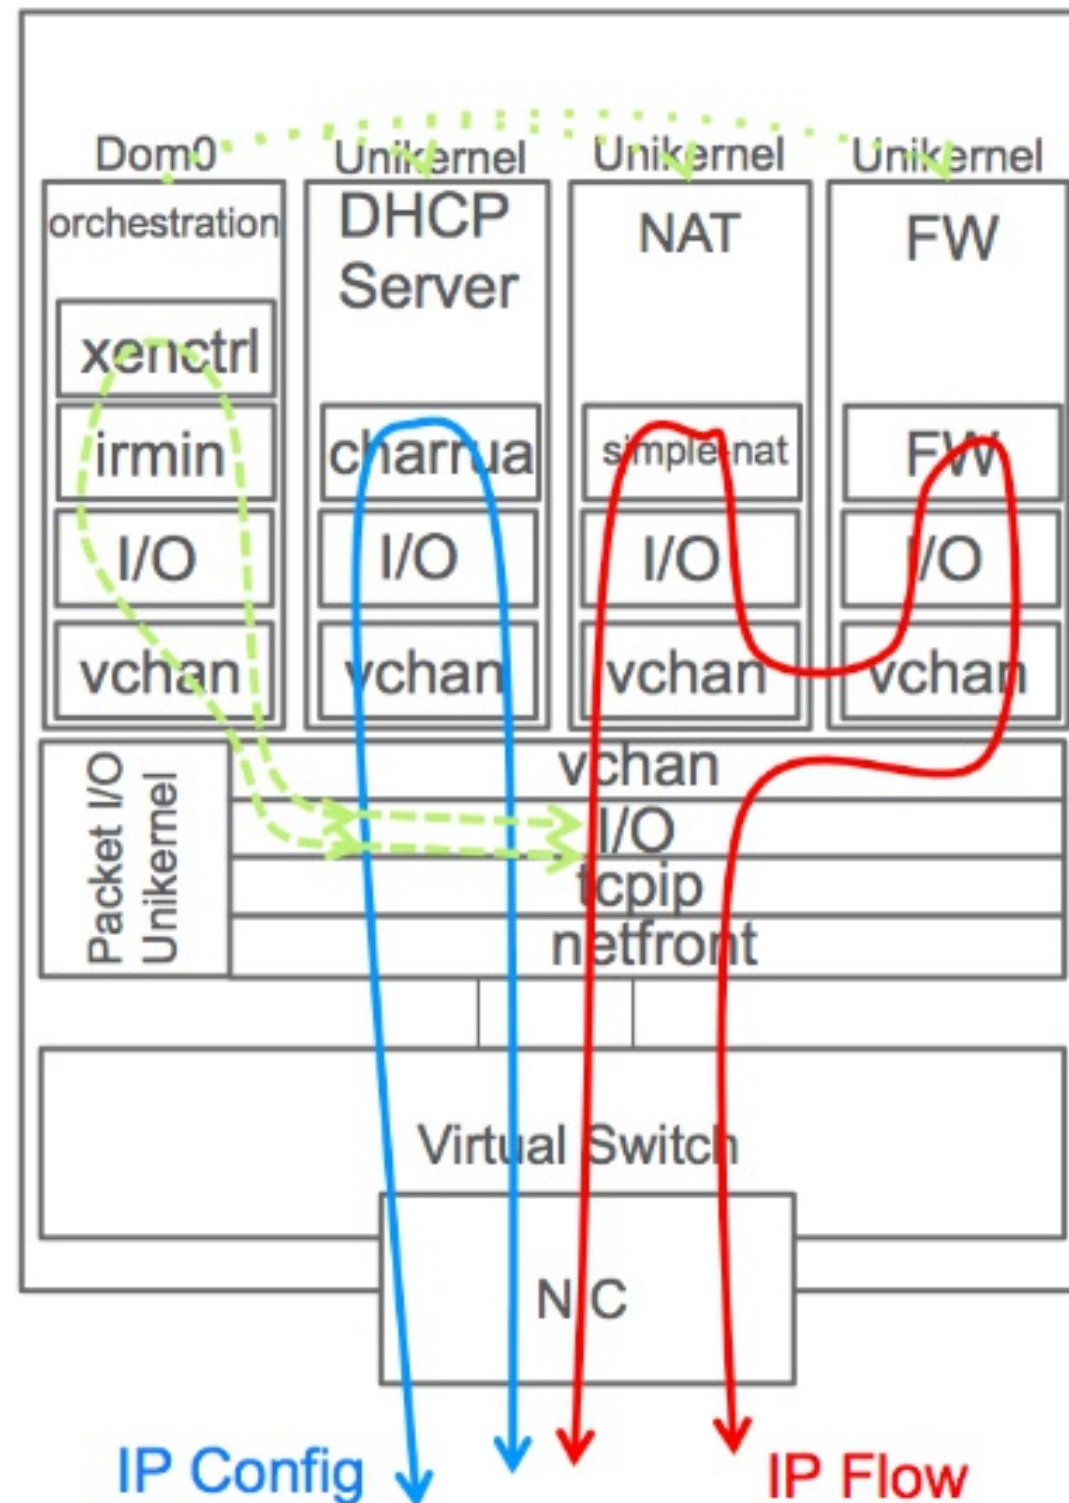
## IoT / Embedded

- For low-resource, potentially secure elements (baremetal or  $\mu$ -vmm ?)
- Build up the "app" instead of stripping down the "OS"

## HPC

@mjg11 • Greater performance possible (but may be hard work)

## In what domains might they be used? - NFV/SDN



# Unikernel Implementations - 2 families

There are 2 main classes of Unikernels

# Unikernel Implementations - 2 families

There are 2 main classes of Unikernels

The Clean-Slate approach emphasizes safety and security.  
Same language for application and Library OS components.

- MirageOS (Ocaml)
- HalVM (Haskell)
- LING (Erlang)



# Unikernel Implementations - 2 families

There are 2 main classes of Unikernels

The Clean-Slate approach emphasizes safety and security.  
Same language for application and Library OS components.

- MirageOS (Ocaml)
- HalVM (Haskell)
- LING (Erlang)

The Legacy approach favours backward compatibility of existing applications based on POSIX-compatibilities.

Many applications have been ported

- OSv (Tomcat, Jetty, Cassandra, OpenJDK, ...)
- Rumprun (MySQL, PHP, Nginx)
- Runtimejs
- Clive (Go)

# Unikernel Implementations

| Technology  | Description  |
|---|--|
| ClickOS<br><a href="http://cnp.neclab.eu">cnp.neclab.eu</a>         | For embedded network h/w.<br>~5MB images, boots <20ms, 45 µs delay, 100 VMs => 10Gbps                                |
| Clive<br><a href="http://lsub.org">lsub.org</a>                     | Written in Go. For distributed and cloud.  |
| Drawbridge<br><a href="#">MS</a>                                    | Research prototype. Picoprocess/container with minimal kernel API surface, and Windows library OS.                   |
| Graphene<br><a href="#">graphene</a>                                | Securing "multi-process" legacy apps - adds IPC.   |
| HaLVM<br><a href="http://galois.com">galois.com</a>                 | Port of GHC (Glasgow Haskell Compiler) suite.<br>Write apps in Haskell to run on Xen.                                |
| IncludeOS<br><a href="http://includeos.org">includeos.org</a>       | Research project for C++ code on virtual hardware.   |
| LING<br><a href="http://erlangonxen.org">erlangonxen.org</a>        | Erlang/OTP runs on Xen.  |
| MirageOS<br><a href="http://mirage.io">mirage.io</a>                | Clean-slate library OS for secure, high-perf network apps.<br>More than 100 MirageOS libraries plus OCaml ecosystem. |
| <a href="#">@mibright</a> OSv<br><a href="http://osv.io">osv.io</a> | Run Linux binaries (w. limitations), supports C/C++, JVM, Ruby, Node.js  |
| Cloudius  |  |
| Rumprun   | FreeBSD. Runs POSIX s/w on BM or VM (Xen)  |

# Unikernel implementations - MirageOS/Ocaml

Clean-Slate

**MIRAGE OS**

<https://mirage.io/>

MirageOS "Library OS" components are written in **Ocaml**.

ML-derived languages are best known for their static type systems and type-inferring compilers.

OCaml-Based



OCaml unifies functional, imperative, and object-oriented programming under an ML-like type system.

OCaml has extensive libraries available

(Unison utility)

# Unikernel implementations - MirageOS-2

Clean-Slate



<https://mirage.io/>

MirageOS Unikernels are based on the Mirage-OS Unikernel base (OS library).

The mirage tool is used to build Unikernels for various backends:

OCaml-Based



- Xen Hypervisor (PV)
- Unix (Linux or OS/X binaries)
- Browser (via Ocaml->JS compiler !!)
- MirageOS 3 (/Solo5) will support kvm (/ukvm)
- Even an experimental BM backend for Raspberry Pi

@mjbright

# Unikernel implementations - MirageOS-2

Clean-Slate



<https://mirage.io/>

OCaml-Based



MirageOS Unikernels are based on the Mirage-OS Unikernel base (OS library).

The mirage tool is used to build Unikernels for various backends:

- Xen Hypervisor (PV)
- Unix (Linux or OS/X binaries)
- Browser (via Ocaml->JS compiler !!)
- MirageOS 3 (/Solo5) will support kvm (/ukvm)
- Even an experimental BM backend for Raspberry Pi

Building applications for unix or xen

```
mirage configure -t unix
make
./mir-console
```

```
mirage configure -t xen
make
****xen create ./mir-console.xen
```


@mjbright

# Unikernel implementations - MirageOS - Use Cases

Clean-Slate

**MIRAGE OS**

<https://mirage.io/>



- BNC Pinata: <http://ownme.ipredator.se/>
- Networking applications
  - e.g. CyberChaff "false network hosts"
- PayGarden, Sean Grove
  - ▶ ["Baby steps to unikernels in production"](#)
    - Too painful to create/configure AMI images on AWS
    - Solo5 allows to create KVM images deployable on GCE

@mjbright

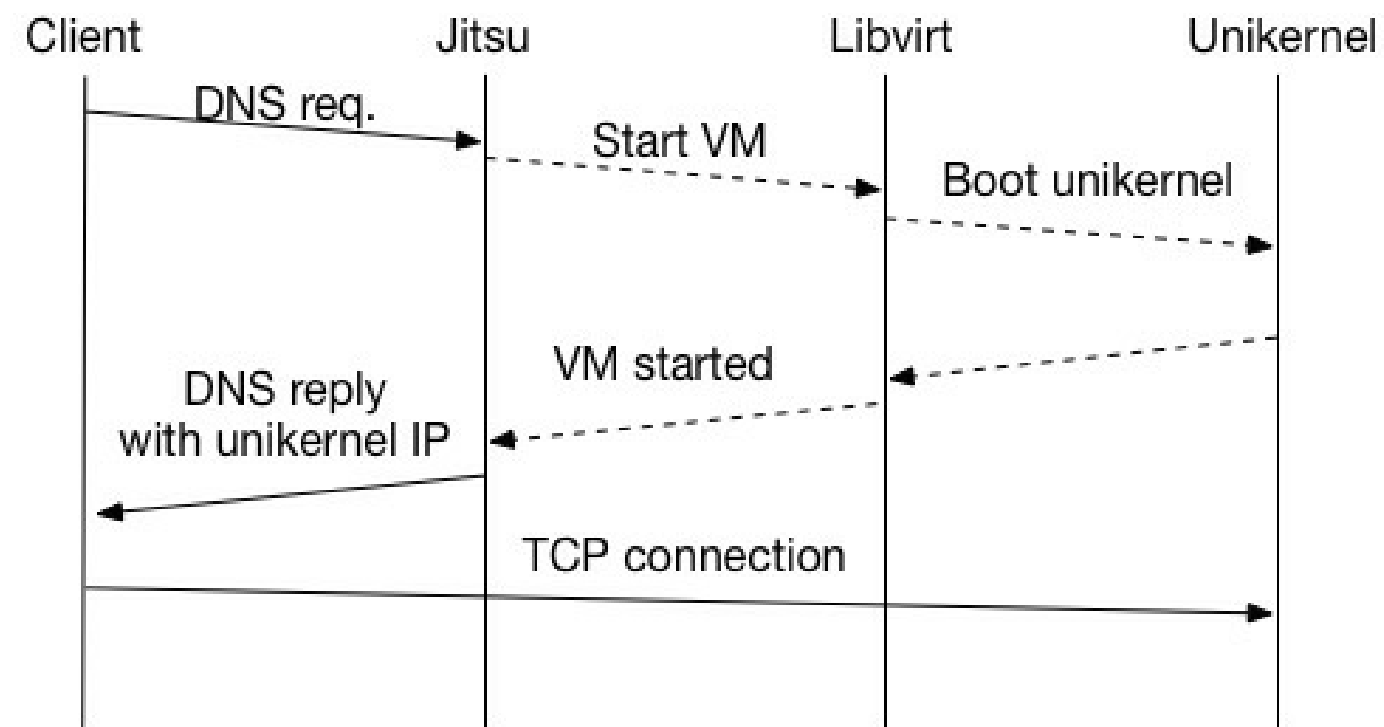
# Unikernel Tooling

- **Unik [EMC-Dell]**: "The Unikernel Compilation and Deployment Platform" (+ image hub)
  - rumprun: Python, Node.js and Go
  - OSv: Java, Node.js, C and C++
  - IncludeOS: C++
  - MirageOS: OCaml
- **Solo5 [IBM]** : An alternative unikernel-base for MirageOS
  - Provides qemu/KVM support for MirageOS
  - Is currently being integrated into MirageOS 3 beta
- **ukvm [IBM]**: An alternative VM Monitor
  - a "library hypervisor"
- **capstan** : OSv build tool (+ image hub)

# Unikernel Tooling

MirageOS jitsu : "Just-In-Time Summoning of Unikernels"

A DNS server that starts unikernels on demand.



Tested with MirageOS and Rumprun unikernels.

<https://github.com/mirage/jitsu>



# Unikernels and Containers : My guess ...

So what about Containers? ... and why did Docker buy Unikernel Systems?

## Unikernels and Containers : My guess ...

So what about Containers? ... and why did Docker buy Unikernel Systems?

- Unikernel Systems are involved in MirageOS/Xen
- Unikernels already used as specific functions in "Docker for Mac"

# Unikernels and Containers : My guess ...

So what about Containers? ... and why did Docker buy Unikernel Systems?

- Unikernel Systems are involved in MirageOS/Xen
- Unikernels already used as specific functions in "Docker for Mac"
- No-brainer: Provide build/ship/run tools for Unikernels
  - build:
    - tools to facilitate building Unikernels
    - test: run Unikernels in containers to facilitate testing  
<https://github.com/mato/docker-unikernel-runner>
  - ship: Docker registry extended to provide Unikernel images
  - run: Docker Swarm orchestrates tasks incl. Unikernels

# Unikernels and Containers : My guess ...

So what about Containers? ... and why did Docker buy Unikernel Systems?

- Unikernel Systems are involved in MirageOS/Xen
- Unikernels already used as specific functions in "Docker for Mac"
- No-brainer: Provide build/ship/run tools for Unikernels
  - build:
    - tools to facilitate building Unikernels
    - test: run Unikernels in containers to facilitate testing  
<https://github.com/mato/docker-unikernel-runner>
  - ship: Docker registry extended to provide Unikernel images
  - run: Docker Swarm orchestrates tasks incl. Unikernels
- Secure Container deployments through hybrid solutions
  - Secure front-ends made of unikernels
    - e.g. for OCaml MediaWiki (http2https, tlstunnel, ...)
  - Containers for backend

@mjrbright  
• Surprises? ...

# Demo



DeferPanic - Unikernel IaaS - <https://deferpanic.com/>

runtime.js - Node.js Unikernel - <https://github.com/runtimejs/example-web-server>

4 unikernel demos - Look Ma, no OS! - <https://github.com/technolo-g/lookma>

# Conclusions



- Much work needs to be done
  - to make them easy to build, deploy, **debug**
  - We will see easier to use solutions
    - Whatever Docker plan to surprise us with
    - Unik will facilitate building, deploying multiple technologies
    - Solo5 will allow mixing of technologies
- Several disparate technologies today
  - but some efforts to synergize
- Unikernels are an interesting complimentary technology to containers
  - We can expect hybrid solutions
- 2017 will be an interesting year for Unikernels

# Resources



Scoop.it  
Unikernels

[www.scoop.it/t/unikernels](http://www.scoop.it/t/unikernels)



Youtube  
Playlist

[youtube.com/.../unikernels](https://youtube.com/.../unikernels)



Wikipedia

[en.wikipedia.org/wiki/Unikernel](http://en.wikipedia.org/wiki/Unikernel)

unikernel


unikernels.org

unikernels.org



mirageos.io

mirageos.io  
[mirage.io/docs/papers](http://mirage.io/docs/papers)



O'Reilly  
"Unikernels"

Free download



@unikernel

@unikernel



[github.com/ocaml-labs](https://github.com/ocaml-labs)

ocaml-labs



[github.com/mirage](https://github.com/mirage)

MirageOS

@mjbright

Thank you

Q&A